

# SafeSim: Neuro-Symbolic Medical Text Simplification

Arunbh Yashaswi, Abhishek Rithik Origanti,  
Matheshwara Annamalai Senthilkumar



# Addressing Hallucination in Medical Text

SafeSim combines LLM fluency with symbolic verification.

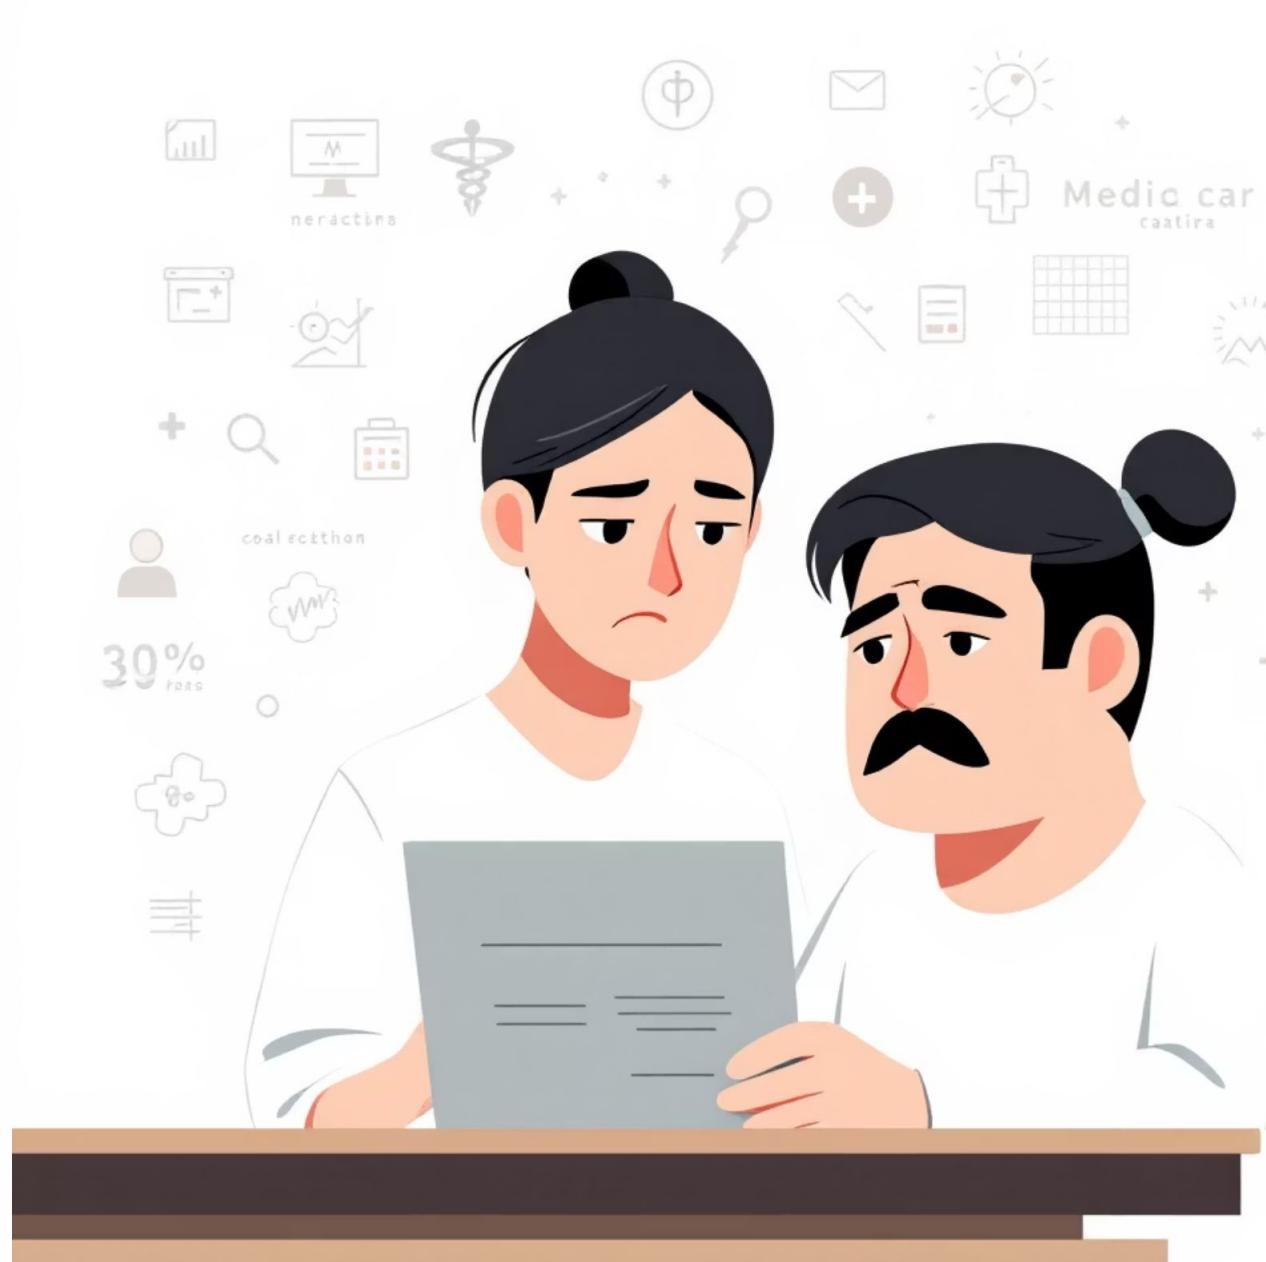
100% Entity Preservation  
Ensures critical medical entities (dosages, vitals) are never altered.

Maintains Readability  
Simplifies complex medical jargon without compromising clarity.

Novel Approach  
Outperforms neural baselines like BART and T5 in safety.

# The Challenge: Patient Understanding

Complex medical jargon creates barriers, leading to poor adherence and adverse outcomes.



- Complex Jargon  
Discharge summaries and clinical notes are often difficult to understand.
- Safety Risks  
LLMs can hallucinate, changing "50mg" to "5mg" with life-threatening consequences.
- Need for Precision  
Automated simplification requires deterministic fact preservation.

# SafeSim: A Neuro-Symbolic Framework

Ensuring deterministic fact preservation without sacrificing readability.

O1

## Symbolic Entity Extraction

Extracts critical entities (dosages, medications, vitals) using symbolic methods.

O

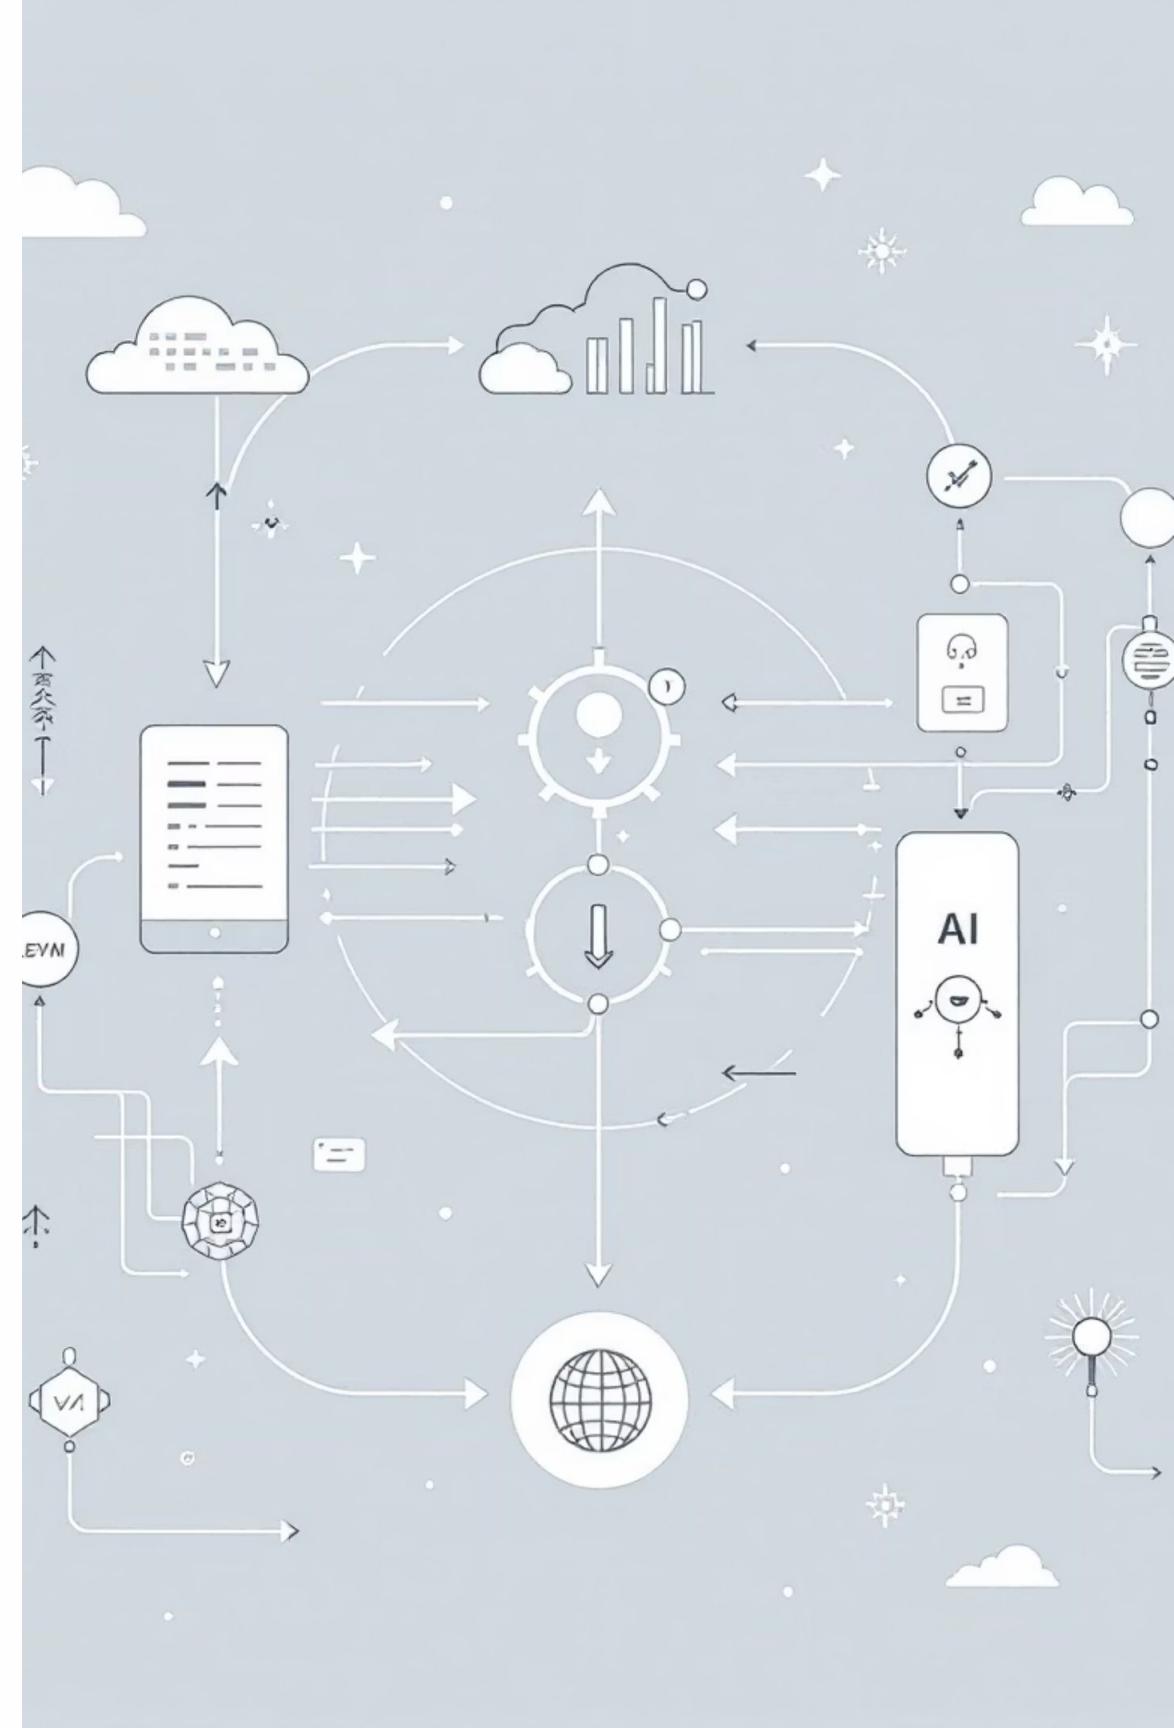
## Neural Simplification

Simplifies text via LLMs, with extracted entities as explicit constraints.

O

## Deterministic Verification

Verifies outputs using logic checks; failed verifications trigger refinement.



# Evaluation Dimensions

SafeSim assessed across five key areas for robust performance.



## System Design

Modular pipeline: NER, LLM  
rewriting, logic verification.



## Method Comparison

SafeSim vs. BART and T5  
baselines for safety and quality.



## Data-Centric NLP

Performance across diverse  
medical specialties and  
terminology.



## Error Analysis

Qualitative audit of failure  
modes, distinguishing critical  
safety violations.



## Responsible NLP

Ethical implications, bias, and  
human-in-the-loop deployment.



# Problem Formulation: Constrained Translation

Maximizing readability while ensuring strict safety constraints.

Source Sequence (S)

Complex medical tokens.

Target Sequence (T)

Simplified, readable text.

Critical Entities (ES)

High-risk categories: Dosage, Medication, Frequency.

Safety Constraint (C)

Semantic invariance between S and T for critical entities.

## Extracted Entities

Give **10 units** **insulin** subcu  
**b.i.d** before meals.

### Entity List:

**10 units** (DOSAGE)

**insulin** (MEDICATION)

**b.i.d** (FREQUENCY)

SafeSim uses a **neuro-symbolic approach** to simplify medical discharge summaries while guaranteeing the preservation of critical facts (medications, dosages, vitals).

## Input: Medical Text

Paste discharge summary or clinical note:

Give 10 units insulin subcutaneously b.i.d  
before meals.

Simplify Text

## Output: Simplified Text

**SAFE**

**Verification Score:** 100%

Simplified text:

Give 10 units insulin under the skin twice a day  
before meals.

Detailed Analysis

## Verification Details

```
▼ {  
  "is_safe" : true
  ▶ "missing_entities"
  ▶ "modified_entities"
  ▶ "warnings" : []
  "score" : 1
}
```

# SafeSim: Safe Medical Text Simplification

SafeSim uses a **neuro-symbolic approach** to simplify medical discharge summaries while guaranteeing the preservation of critical facts (medications, dosages, vitals).

## Input: Medical Text

Paste discharge summary or clinical note:

Administer 10 units insulin subcutaneously b.i.d. before meals. Check blood glucose q.i.d. Target range 80-120 mg/dL.



## Output: Simplified Text

 NEEDS REVIEW

Verification Score: 80%

Simplified text:

give 10 units insulin under the skin twice a day. before meals. Check blood glucose four times a day. Target range 80-120 mg/dL.

## ⚠ Safety Alerts

⚠ SAFETY ALERT: Critical entity 'Administer' not found in simplified text!

### Detailed Analysis

## Extracted Entities

Administer    10 units    insulin

subcutaneously    b.i.d.    before meals. Check  
blood glucose    q.i.d.    Target    range 80-  
120 mg /dL.

## Entity List:

Administer (MEDICATION)

10 units (DOSAGE)

insulin (MEDICATION)

b.i.d. (FREQUENCY)

q.i.d. (FREQUENCY)

## Verification Details

```
▼ {  
  "is_safe" : true  
  ▶ "missing_entities" : []  
  ▶ "modified_entities" : []  
  ▶ "warnings" : []  
  "score" : 1  
}
```

## Model Information

Backend: rule-based

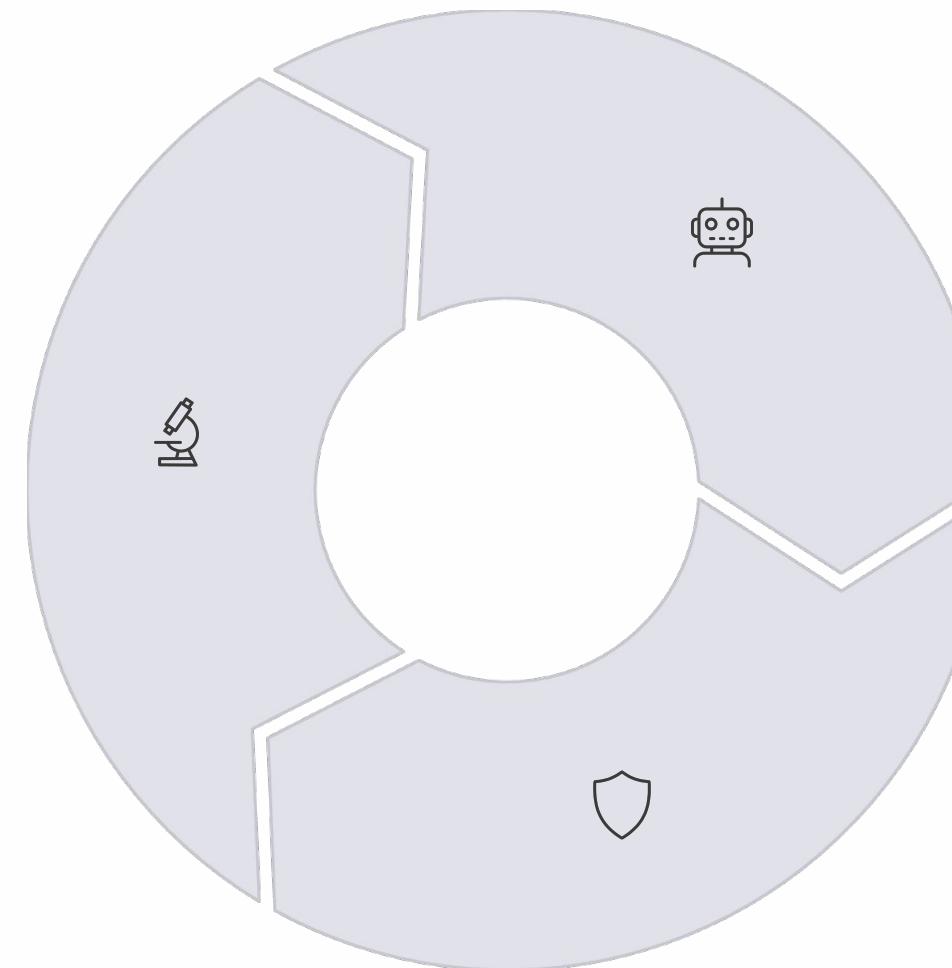
Strictness: high

Try SafeSim Demo



# SafeSim Architecture: Verify and Refine

A neuro-symbolic pipeline for robust medical text simplification.

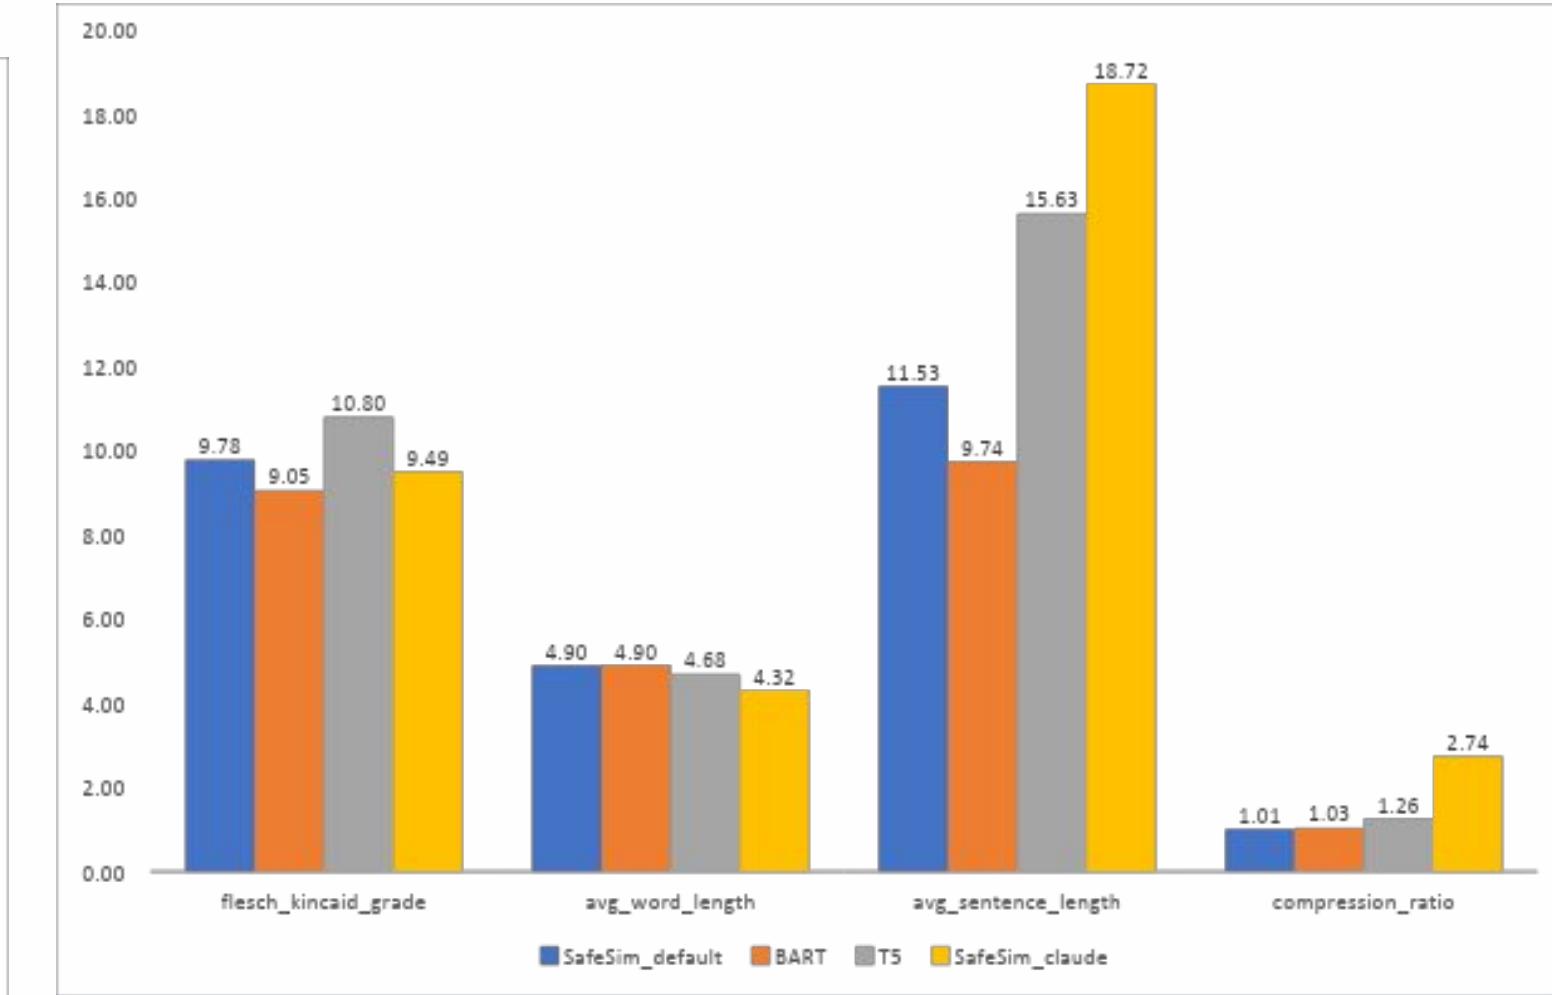
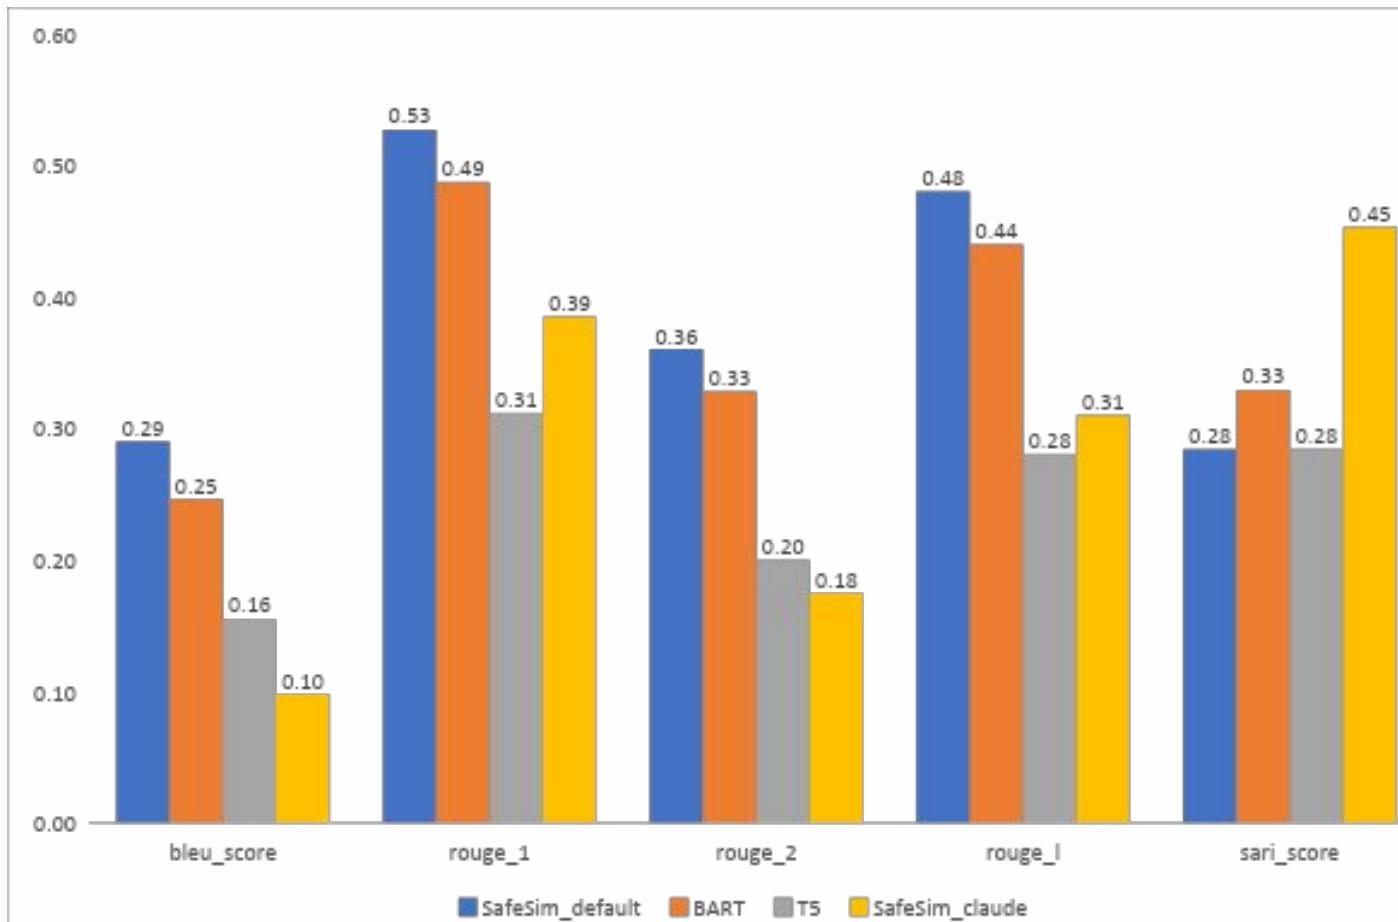


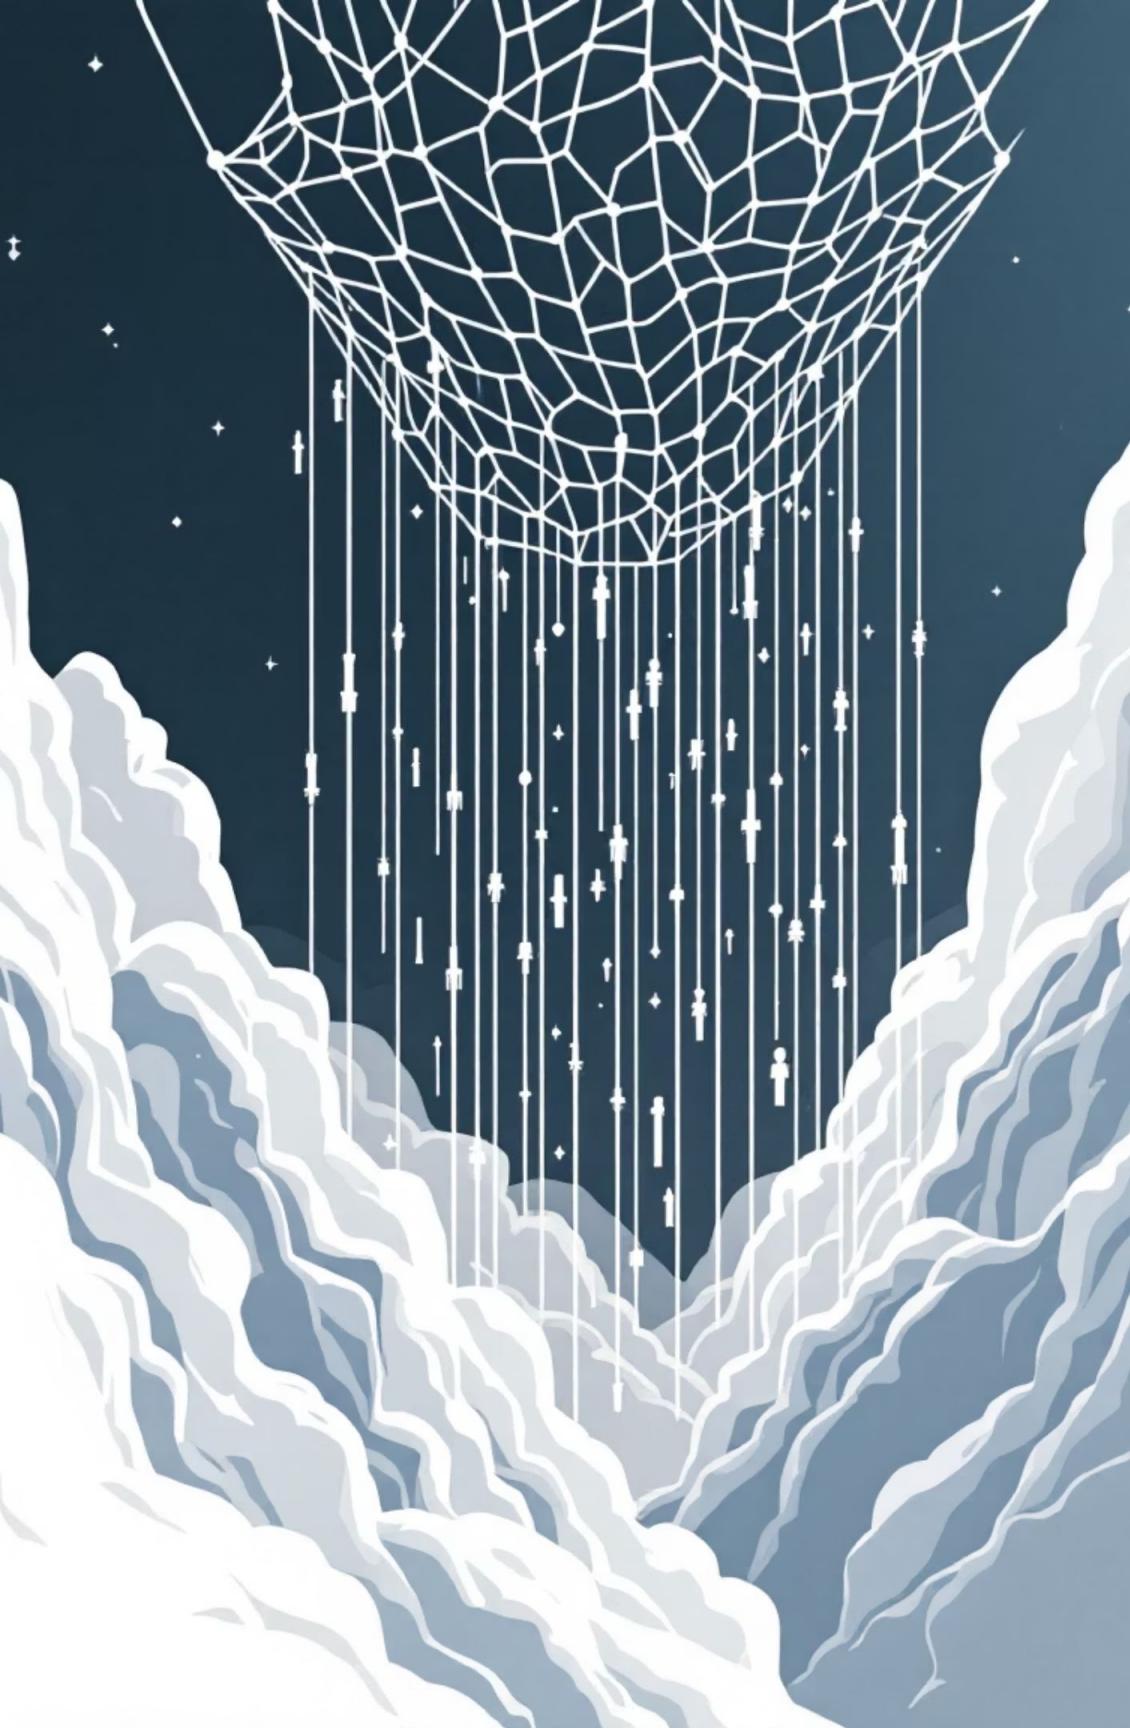
Symbolic Entity Extraction  
Scispacy NER + Regex for structured  
and unstructured entities.

Neural Simplification  
LLM backend (GPT-4, Claude, BART)  
with entity injection.

Deterministic Verification  
Re-extract entities, set comparison,  
semantic relaxation, regeneration on  
failure.

# Performance Metrics: SafeSim vs. Baselines





# Key Findings & Discussion

SafeSim's neuro-symbolic approach ensures safety and effective simplification.



Simplicity-Overlap  
Trade-off  
High SARI, low BLEU indicates  
aggressive, effective simplification,  
not just copying.



Neuro-Symbolic Safety  
100% Entity & Dosage Preservation,  
eliminating stochastic risk of pure  
neural models.



Structural Readability  
Merges fragmented shorthand into coherent sentences, improving patient  
comprehension.



# Conclusion & Future

## Work

SalesSim solves hallucination in medical text simplification, setting a new standard.

### Key Achievements

- 100% Entity Preservation Rate (EPR)
- Clinically safe LLM-based simplification
- Effective rewriting, not just text copying

### Future Directions

- Expand entity coverage (e.g., conditional dosing)
- Multilingual support for global accessibility
- End-to-end optimization with RL feedback

**THANK YOU**